

Probabilistic Machine Learning

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN



Exercise 7

Summer Term 2023

Full Name	Matriculation No
Batuhan Özçömlekci	6300476
Aakarsh Nair	6546577

1. Theory Question : This theory exercise is about the properties of block matrices. That is matrices of the form

$$M = \begin{bmatrix} A & B \\ C & D \end{bmatrix}, \text{ where } A \text{ and } D \text{ are square blocks of arbitrary size.}$$

You can assume all the inverses below exist.

(a) Inverse of a block matrix: In the lecture, the following statement about inverse of M was provided without proof.

$$M^{-1} = \begin{bmatrix} A^{-1} + \bar{A}^{-1} B Q^{-1} C A^{-1} & -\bar{A}^{-1} B Q^{-1} \\ -Q^{-1} C A^{-1} & Q^{-1} \end{bmatrix}$$

$Q := D - C A^{-1} B$, denotes the so called Schur complement. Prove this statement.

Answer :- We construct a set of reductions E_1, E_2, E_3, E_4 such that

$$(E_4 E_3 E_2 E_1) M = I \Rightarrow M^{-1} = (E_4 E_3 E_2 E_1)$$

$$\textcircled{1} \text{ Let } E_1 = \begin{bmatrix} A^{-1} & 0 \\ 0 & I \end{bmatrix} \Rightarrow E_1 M = \begin{bmatrix} A^{-1} & 0 \\ 0 & I \end{bmatrix} \begin{bmatrix} A & B \\ C & D \end{bmatrix} = \begin{bmatrix} I & A^{-1} B \\ C & D \end{bmatrix} := R_1$$

$$\textcircled{2} \text{ Let } E_2 = \begin{bmatrix} I & 0 \\ -C & I \end{bmatrix} \Rightarrow E_2 R_1 = \begin{bmatrix} I & 0 \\ -C & I \end{bmatrix} \begin{bmatrix} I & A^{-1} B \\ C & D \end{bmatrix} = \begin{bmatrix} I & A^{-1} B \\ C-C & D-CA^{-1}B \end{bmatrix} := R_2$$

$$\textcircled{3} \text{ Let } E_3 = \begin{bmatrix} I & 0 \\ 0 & (D-CA^{-1}B)^{-1} \end{bmatrix} \Rightarrow E_3 R_2 = \begin{bmatrix} I & 0 \\ 0 & (D-CA^{-1}B)^{-1} \end{bmatrix} \begin{bmatrix} I & A^{-1} B \\ C-C & D-CA^{-1}B \end{bmatrix} \\ = \begin{bmatrix} I & A^{-1} B \\ 0 & I \end{bmatrix} := R_3$$

$$\textcircled{4} \text{ Let } E_4 = \begin{bmatrix} I & -A^{-1}B \\ 0 & I \end{bmatrix} \Rightarrow E_4 R_3 = \begin{bmatrix} I & -A^{-1}B \\ 0 & I \end{bmatrix} \begin{bmatrix} I & A^{-1} B \\ 0 & I \end{bmatrix} = \begin{bmatrix} I & 0 \\ 0 & I \end{bmatrix}$$

$$\begin{aligned}
M^{-1} &= \begin{pmatrix} E_4 & E_3 & E_2 & E_1 \end{pmatrix} = \begin{bmatrix} I & -A'^{-1}B \\ 0 & I \end{bmatrix} \begin{bmatrix} I & 0 \\ 0 & (D - CA'^{-1}B)^{-1} \end{bmatrix} \begin{bmatrix} I & 0 \\ -C & I \end{bmatrix} \begin{bmatrix} A' & 0 \\ 0 & I \end{bmatrix} \\
&= \begin{bmatrix} I & -A'^{-1}B(D - CA'^{-1}B)^{-1} \\ 0 & (D - CA'^{-1}B)^{-1} \end{bmatrix} \begin{bmatrix} A' & 0 \\ -CA'^{-1} & I \end{bmatrix} \\
&= \begin{bmatrix} A'^{-1} + A'^{-1}B(D - CA'^{-1}B)^{-1}CA'^{-1} & -A'^{-1}B(D - CA'^{-1}B)^{-1} \\ -(D - CA'^{-1}B)^{-1}CA'^{-1} & (D - CA'^{-1}B)^{-1} \end{bmatrix} \\
&= \begin{bmatrix} A'^{-1} + A'^{-1}BQ^{-1}CA'^{-1} & -A'^{-1}BQ^{-1} \\ -Q^{-1}CA'^{-1} & Q^{-1} \end{bmatrix}
\end{aligned}$$

■

(b) Block LU decomposition: The LU decomposition of M is a decomposition into a lower triangular matrix L and upper triangular matrix U such that $M = LU$

Show that LU decomposition of M is given by.

$$L = \begin{bmatrix} A^{1/2} & 0 \\ C(A^{1/2})^T & Q^{1/2} \end{bmatrix} \text{ and } U = \begin{bmatrix} (A^{1/2})^T & A^{-1/2}B \\ 0 & (Q^{1/2})^T \end{bmatrix}$$

Answer: Forward multiplying LU we get

$$\begin{aligned} LU &= \begin{bmatrix} A^{1/2} & 0 \\ C(A^{1/2})^T & Q^{1/2} \end{bmatrix} \begin{bmatrix} (A^{1/2})^T & A^{-1/2}B \\ 0 & (Q^{1/2})^T \end{bmatrix} \\ &= \begin{bmatrix} A^{1/2}(A^{1/2})^T & A^{1/2}A^{-1/2}B \\ C(A^{1/2})(A^{1/2})^T & C(A^{1/2})^T A^{-1/2}B + Q^{1/2}Q^{1/2}^T \end{bmatrix} \\ &= \begin{bmatrix} A & B \\ C & D - C\bar{A}^T B + CA^T B \end{bmatrix} = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \end{aligned}$$

Using
 ① $A^{1/2}A^{-1/2} = I$
 ② $Q^{1/2}Q^{1/2}^T = I$

We can also derive the expression for LU using block elimination.

$$\text{Let } E_1 = \begin{bmatrix} A^{-1/2} & 0 \\ 0 & I \end{bmatrix}$$

$$\text{Let } f_2 = \begin{bmatrix} I & 0 \\ -C\bar{A}^{1/2} & I \end{bmatrix}$$

$$E_1 M = \begin{bmatrix} A^{-1/2} & 0 \\ 0 & I \end{bmatrix} \begin{bmatrix} A & B \\ C & D \end{bmatrix} = \begin{bmatrix} A^{-1/2}^T & \bar{A}^{-1/2}B \\ C & D \end{bmatrix} := R_1$$

$$E_1 R_1 = \begin{bmatrix} I & 0 \\ -C\bar{A}^{1/2} & I \end{bmatrix} \begin{bmatrix} A^{-1/2}^T & \bar{A}^{-1/2}B \\ C & D \end{bmatrix} = \begin{bmatrix} A^{-1/2}^T & \bar{A}^{-1/2}B \\ -C\bar{A}^{1/2}A^{-1/2}^T + C & D - C\bar{A}^{1/2}\bar{A}^{-1/2}B \end{bmatrix}$$

$$= \begin{bmatrix} A^{-1/2}^T & \bar{A}^{-1/2}B \\ 0 & D - C\bar{A}^{1/2}B \end{bmatrix} := R_2$$

$$\text{Let } E_3 = \begin{bmatrix} I & 0 \\ 0 & (D - CA^T B)^{-1/2} \end{bmatrix} \Rightarrow E_3 R_2 = \begin{bmatrix} I & 0 \\ 0 & \underbrace{(D - CA^T B)^{-1/2}}_{Q^{1/2}} \end{bmatrix} \begin{bmatrix} A^{1/2} \\ 0 \end{bmatrix}$$

$$= \begin{bmatrix} A^{1/2} & A^{1/2} B \\ 0 & (D - CA^T B)^{1/2} \end{bmatrix}$$

$$\text{Thus } (E_3 E_2 E_1) M = U$$

$$\Rightarrow M = (E_3 E_2 E_1)^{-1} U = E_1^{-1} E_2^{-1} E_3^{-1} U$$

$$E_1^{-1} = \begin{bmatrix} A^{1/2} & 0 \\ 0 & I \end{bmatrix} \quad E_2^{-1} = \begin{bmatrix} I & 0 \\ C A^{-1/2} & I \end{bmatrix} \quad E_3^{-1} = \begin{bmatrix} I & 0 \\ 0 & Q^{1/2} \end{bmatrix}$$

using block diagonal inverse formula above.

$$E_1^{-1} E_2^{-1} E_3^{-1} = \begin{bmatrix} A^{1/2} & 0 \\ 0 & I \end{bmatrix} \begin{bmatrix} I & 0 \\ C A^{-1/2} & I \end{bmatrix} \begin{bmatrix} I & 0 \\ 0 & Q^{1/2} \end{bmatrix}$$

$$= \begin{bmatrix} A^{1/2} & 0 \\ C A^{-1/2} & I \end{bmatrix} \begin{bmatrix} I & 0 \\ 0 & Q^{1/2} \end{bmatrix} = \begin{bmatrix} A^{1/2} & 0 \\ C A^{-1/2} & Q^{1/2} \end{bmatrix}$$

which is the lower triangular matrix as expected.

$$\Rightarrow M = LU \text{ where } L = \begin{bmatrix} A^{1/2} & 0 \\ C A^{-1/2} & Q^{1/2} \end{bmatrix} \quad U = \begin{bmatrix} (A^{1/2})^T & A^{1/2} B \\ 0 & Q^{1/2} \end{bmatrix}$$



1(c) Determinant of a block matrix: Show that

$$\det(M) = \det(P) \det(Q)$$

Hint: You can use that

$$\det \begin{pmatrix} P & Q \\ 0 & S \end{pmatrix} = \det \begin{pmatrix} P & 0 \\ R & S \end{pmatrix} = \det(P) \det(S)$$

and multiplicativity of the determinant $\det(P \cdot Q) = \det(P) \det(Q)$

Answer:

We know $M = LU$ is the LU decomposition for M .

$$\det(M) = \det(LU) = \det(L) \det(U)$$

$$\Rightarrow \det \left(\begin{bmatrix} A^{1/2} & 0 \\ CA^{1/2} & Q^{1/2} \end{bmatrix} \right) \det \left(\begin{bmatrix} A^{1/2} & A^{1/2}B \\ 0 & Q^{1/2} \end{bmatrix} \right) \quad \text{From 1(b) we know}$$

$$\Rightarrow \det(A^{1/2}) \det(Q^{1/2}) \det(A^{1/2}B) \det(Q^{1/2})$$

$$\Rightarrow \det(A^{1/2}A^{1/2}) \det(Q^{1/2}Q^{1/2})$$

$$= \det(A) \det(Q)$$



Exercise 07

June 19, 2023

1 Probabilistic Machine Learning

University of Tübingen, Summer Term 2023 © 2023 P. Hennig

1.1 Exercise Sheet No. 7 — Online Gaussian Process Regression

Submission by:

- Batuhan, Oezcoemlekci, Matrikelnummer: 6300476
- Aakarsh, Nair, Matrikelnummer: 6546577

```
[ ]: # For measuring wallclock time
from time import time

# Numpy, JAX
import jax
import jax.numpy as jnp
import numpy as np
jax.config.update("jax_enable_x64", True) # use double-precision numbers
```

```
[ ]: # Plotting
from matplotlib import pyplot as plt
from tueplots import bundles

import matplotlib as mpl

%config InlineBackend.figure_formats = ["svg"]
%matplotlib inline
plt.rcParams.update(bundles.beamer_moml())
plt.rcParams.update({'figure.dpi': 200})
```

```
[ ]: # Import `gaussians.py`
import gaussians
```

1.1.1 0. Overview

In this exercise sheet, we are going to implement Gaussian Process (GP) regression in an *online* setting. That means, data arrives sequentially over time and the GP regression model is incrementally updated to incorporate that new information.

In the lecture and previous exercise sheets, you have already seen and worked with the classes `Gaussian`, `GaussianProcess`, and `ConditionalGaussianProcess` implemented in `gaussians.py`. This functionality is also very straightforward to use in the online setting, but a naive approach will turn out to be prohibitively inefficient. We are going to investigate the underlying problem in more detail and then develop and implement a more efficient strategy by using some tools from linear algebra (like the Schur complement and a block-wise Cholesky factorization).

1.1.2 1. Training and Test Data

Since we want to simulate that the training data is provided *incrementally*, it is represented by two lists `X_list` (containing 6 subsets of 100 data points each) and `Y_list` (containing the corresponding regression targets). As a first step, we load the data and visualize these subsets.

```
[ ]: # Load data
data = np.load("data.npz")

# Extract from data
X_list = data["X_list"]
Y_list = data["Y_list"]
sigma = data["sigma"]
X_test = data["X_test"]
```

```
[ ]: # Check shape of sections
print("Training data")
for idx, (Xi, Yi) in enumerate(zip(X_list, Y_list)):
    print(f"Subset index {idx}: X.shape = {Xi.shape}, Y.shape = {Yi.shape}")

print("\nTest data")
print("X_test.shape = ", X_test.shape)
```

```
Training data
Subset index 0: X.shape = (100, 1), Y.shape = (100,)
Subset index 1: X.shape = (100, 1), Y.shape = (100,)
Subset index 2: X.shape = (100, 1), Y.shape = (100,)
Subset index 3: X.shape = (100, 1), Y.shape = (100,)
Subset index 4: X.shape = (100, 1), Y.shape = (100,)
Subset index 5: X.shape = (100, 1), Y.shape = (100,)
```

```
Test data
X_test.shape = (100, 1)
```

```
[ ]: def setup_plot(ax, title=None):
    ax.set_xlabel("$x$")
```

```

ax.set_ylabel("$y$")
ax.grid()
if title is not None:
    ax.set_title(title)

# Visualization of training data
fig, ax = plt.subplots(1, 1, figsize=(4, 3))

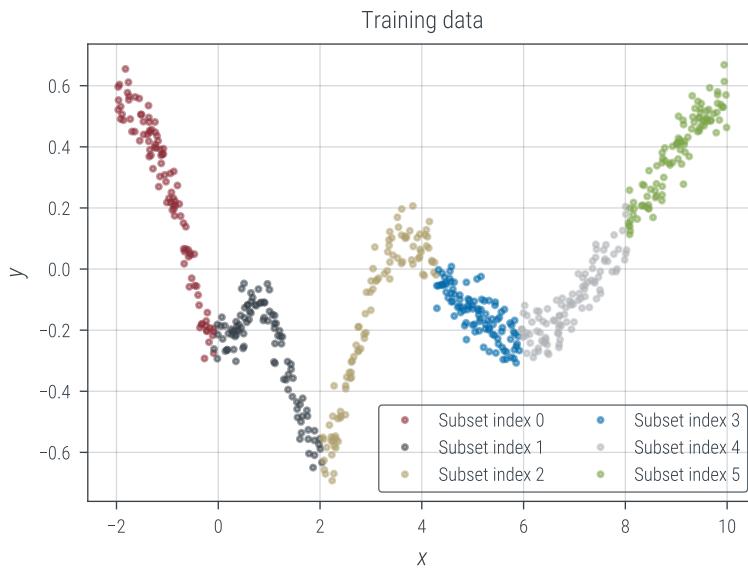
for idx, (Xi, Yi) in enumerate(zip(X_list, Y_list)):
    ax.plot(Xi, Yi, "o", ms=2, alpha=0.5, label=f"Subset index {idx}")
print(Xi.shape)

setup_plot(ax, "Training data")
ax.legend(ncol=2)

plt.show()

```

(100, 1)



1.1.3 2. Prior GP

Here, we set up the prior GP with a constant mean at zero and rational quadratic kernel as covariance function.

[]: `import functools`

```
def constant_mean(x, c=0.0):
```

```

    return c * jnp.ones_like(x[:, 0])

def RQ_kernel(a, b, ell=1.0, alpha=1.0, theta=1.0):
    return theta**2 * (1 + jnp.sum((a - b) / ell, axis=-1) ** 2 / (2 * alpha)) ↴
    ↴** (-alpha)

```

```
[ ]: # Define mean, kernel and prior GP
mean = functools.partial(constant_mean, c=0.0)
kernel = functools.partial(RQ_kernel, ell=1.0, alpha=1.0, theta=0.5)
prior = gaussians.GaussianProcess(mean, kernel)
```

1.1.4 3. Approach 1

We can use the functionality implemented in `gaussians.py` to perform inference in an incremental way: In the first step, we condition the prior on the first subset (`X_list[0]`, `y_list[0]`). From this, we obtain an instance of a `ConditionalGaussianProcess` that we can condition on the next subset (`X_list[1]`, `y_list[1]`) etc.

```
[ ]: def get_posteriors():
    """Construct posteriors from prior by conditioning on subsets"""

    posteriors = [prior]
    for posterior_idx, (Xi, Yi) in enumerate(zip(X_list, Y_list)):

        # Compute posterior
        posterior = posteriors[-1].condition(Yi, Xi, sigma)
        posteriors.append(posterior)
    posteriors.pop(0) # remove prior from list
    return posteriors

posteriors_A1 = get_posteriors()
```

No GPU/TPU found, falling back to CPU. (Set `TF_CPP_MIN_LOG_LEVEL=0` and `rerun` for more info.)

Let's visualize those posteriors.

```
[ ]: from math import ceil

def get_nrows_ncols(nplots):
    nrows = ceil((1/3) * nplots)
    ncols = ceil(nplots / nrows)
    return nrows, ncols

def plot_conditional_gps(list_of_conditional_gps):
```

```
"""Plot a list of conditional GPs in one figure."""
```

```
# Concatenate training data
X = jnp.concatenate(X_list, axis=0)
Y = jnp.concatenate(Y_list, axis=0)

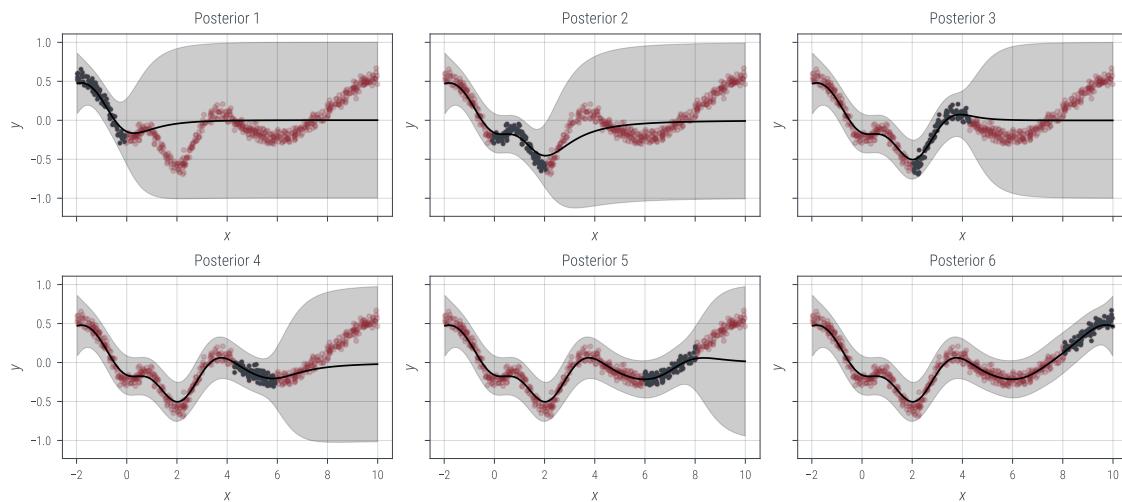
# Define grid of subplots
num_gps = len(list_of_conditional_gps)
nrows, ncols = get_nrows_ncols(num_gps)
fig, axs = plt.subplots(
    nrows,
    ncols,
    figsize=(ncols * 3, nrows * 2),
    sharex=True,
    sharey=True,
)

# Plot GPs
for idx, gp in enumerate(list_of_conditional_gps):
    ax = axs.flatten()[idx]

    # Plot GP and data
    ax.plot(X, Y, ".", ms=5, alpha=0.2)
    ax.plot(gp.X, gp.y, ".", ms=3)
    gp.plot(ax, X_test, color="k", std_kwarg={"alpha": 0.2})

    setup_plot(ax, f"Posterior {idx + 1}")
plt.show()
```

```
plot_conditional_gps(posteriors_A1)
```



Hm, that took quite some time. Let's investigate the runtime in more detail.

```
[ ]: def measure_eval_times(nreps=3):
    """Construct posteriors and measure the time required for evaluating
    them on `X_test`"""

    posteriors = [prior]
    eval_times = np.zeros((nreps, len(X_list)))
    for posterior_idx, (Xi, Yi) in enumerate(zip(X_list, Y_list)):
        for rep_idx in range(nreps):

            # Measure the time for evaluating the posterior on `X_test`
            start_time = time()
            posterior = posteriors[-1].condition(Yi, Xi, sigma)
            posterior(X_test)
            eval_times[rep_idx, posterior_idx] = time() - start_time
            posteriors.append(posterior)

    return eval_times

eval_times_A1 = measure_eval_times()
```

```
[ ]: import matplotlib.colors as mcolors
COLS = ["red", "green", "blue"]

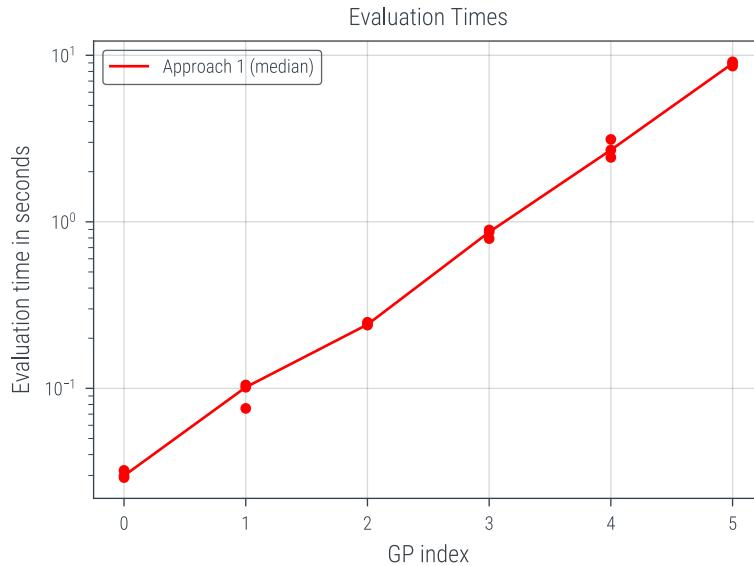
def plot_eval_times(list_of_eval_times):
    """Plot the measured evaluation times. Entry 0 in `list_of_eval_times` is assumed to belong to approach 1, entry 1 to approach 2 etc.
    """
    fig, ax = plt.subplots(1, 1, figsize=(4, 3))

    for idx, eval_times in enumerate(list_of_eval_times):
        ax.plot(eval_times.T, "o", ms=3, c=COLS[idx])
        label = f"Approach {idx + 1} (median)"
        ax.plot(np.median(eval_times, axis=0), "--", c=COLS[idx], label=label)

    # General plot settings
    ax.set_xlabel("GP index")
    ax.set_ylabel("Evaluation time in seconds")
    ax.grid()
    ax.legend()
    ax.set_yscale("log")
    ax.set_title(f"Evaluation Times")

    plt.show()
```

```
plot_eval_times([eval_times_A1])
```



```
[ ]: # Determine the slopes of the log evaluation time
eval_times_A1_median = np.median(eval_times_A1, axis=0)
slopes = eval_times_A1_median[1:] / eval_times_A1_median[:5]
print("slopes = ", slopes)
```

```
slopes = [3.3963706 2.38149446 3.58903933 3.11187118 3.30786873]
```

Observations: The logarithmic plot shows that the runtime increases approximately exponentially reaching up to 10 seconds (on your machine, it may be different of course) for the last posterior. The runtime increases by a factor of around 3-4 for each additional subset of data.

Task: What's going on here? Why is the runtime increasing approximately exponentially?

Your answer: As the data size that the Conditional Gaussian conditioned on increases, the cost of operations (particularly matrix inversions) applied on the current (say N sized) matrix increases as well. Nearly exponential increase is due to matrix inversions (kernel size in Cholesky).

1.1.5 4. Approach 2

Here, we condition the prior on a growing data set. For this, we overwrite the `condition` method of the `ConditionalGaussianProcess` class.

```
[ ]: from gaussians import Gaussian, ConditionalGaussianProcess

def condition_on_prior(self, y_new, X_new, sigma):
    # Check data
```

```

y_new = jnp.atleast_1d(y_new)
X_new = jnp.atleast_2d(X_new)

# Concatenate data
X_both = jnp.concatenate([self.X, X_new], axis=0)
y_both = jnp.concatenate([self.y, y_new], axis=0)

# Concatenate Gaussians
mu_new = jnp.zeros_like(y_new)
mu_both = jnp.concatenate([self.epsilon.mu, mu_new], axis=0)

Sigma_new = sigma * jnp.eye(len(y_new))
Sigma_both = jax.scipy.linalg.block_diag(self.epsilon.Sigma, Sigma_new)

epsilon_both = Gaussian(mu=mu_both, Sigma=Sigma_both)

# Build conditional GP from prior (!)
return ConditionalGaussianProcess(self.prior, y=y_both, X=X_both, ↴
                                     epsilon=epsilon_both)

# Assign method to class
ConditionalGaussianProcess.condition = condition_on_prior

```

```
[ ]: # Evaluate approach 2
posterior_A2 = get_posteriors()
eval_times_A2 = measure_eval_times()
```

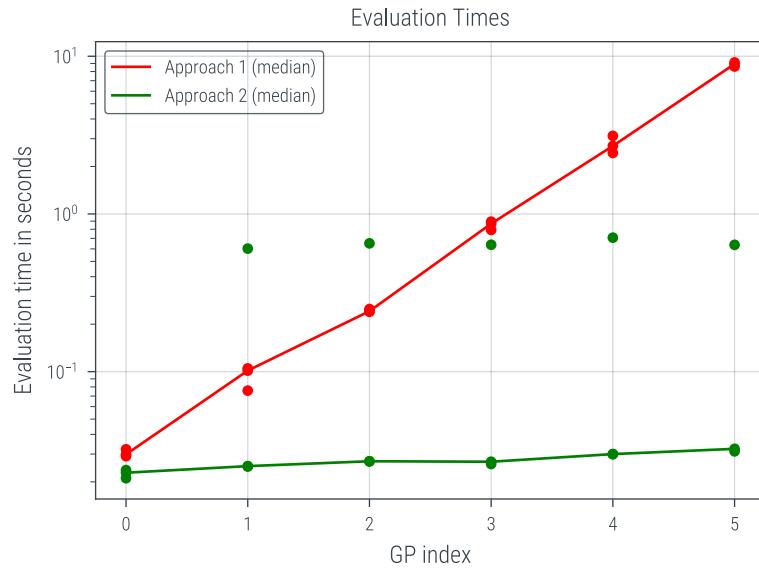
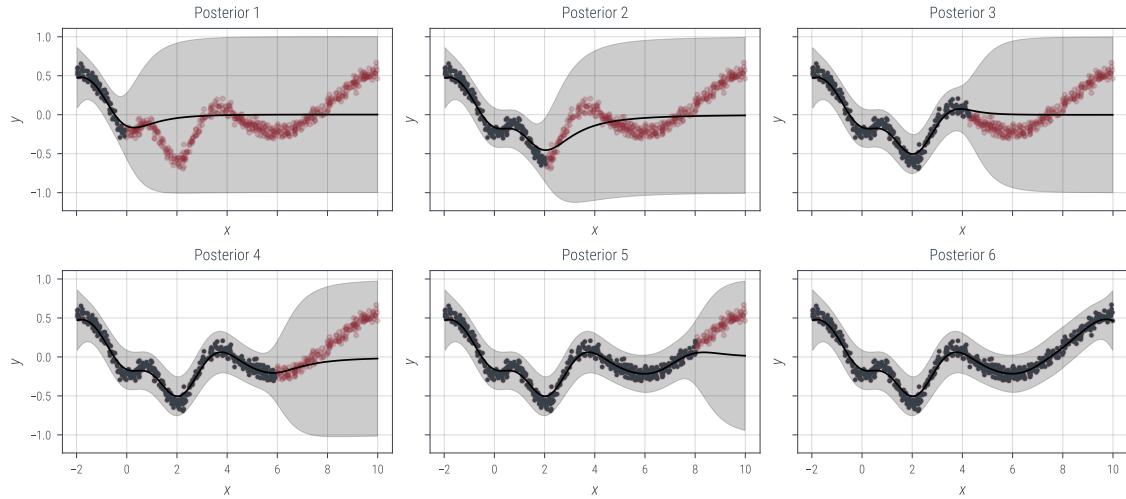
Note: In theory, you can show that the posteriors from approaches 1 and 2 are identical, i.e. conditioning multiple times in a convoluted manner (approach 1) is equivalent to conditioning on the union of the subsets (approach 2). Let's check if we can confirm this claim empirically.

Task: Evaluate the posteriors on `X_test` and check if the resulting Gaussian distributions are identical (`jnp.allclose` might be useful here).

```
[ ]: # TODO
# Check if the resulting Gaussian distributions are identical (`jnp.allclose` ↴
# might be useful here)
for posterior_A1, posterior_A2 in zip(posteriors_A1, posteriors_A2):
    same_mu = jnp.allclose(posterior_A1(X_test).mu, posterior_A2(X_test).mu)
    same_Sigma = jnp.allclose(posterior_A1(X_test).Sigma, posterior_A2(X_test).Sigma)
    if not same_mu or not same_Sigma:
        print("The posteriors are not identical.")
        break
if same_mu and same_Sigma:
    print("The posteriors are identical.")
```

The posteriors are identical.

```
[ ]: # Visualizations
plot_conditional_gps(posteriors_A2)
plot_eval_times([eval_times_A1, eval_times_A2])
```



Even though the slope is less in the plot of approach 2 in logarithmic scale, the (almost) linear trend in the logarithmic scale is present in the second approach too. Empirically, there is a little discrepancy yet the general trend is the same.

1.1.6 5. Approach 3

While approach 2 is less inefficient, it is still quite wasteful: For every new subset, the posterior GP is built entirely from scratch! This can be improved by exchanging information between the

posteriors: Simply put, a posterior that is conditioned on some data set should inform the next posterior what it has already learned about the data. Before we dive into that, we need to get the theory straight.

Our goal is to condition the GP on a data set (X, y) that is the union of (X_{old}, y_{old}) and (X_{new}, y_{new}) . Let's assume that we have conditioned a GP on (X_{old}, y_{old}) and have already triggered the computation of its properties. So, in particular, the following quantities are already cached:

1. **predictive_covariance**: $k(X_{old}, X_{old}) + \sigma^2 I$
2. **predictive_mean**: $m(X_{old}) + \mu(X_{old})$
3. **predictive_covariance_cho**: the Cholesky factorization of **predictive_covariance**
4. **representer_weights**: $[k(X_{old}, X_{old}) + \sigma^2 I]^{-1}(y_{old} - m(X_{old}) + \mu(X_{old}))$

Improved version: Now, the core idea is to use the quantities above to save computations when conditioning on (X, y) .

5. **predictive_covariance**: This matrix now has the following block structure:

$$k(X, X) + \sigma^2 I = \begin{array}{c|c} k(X_{old}, X_{old}) + \sigma^2 I & k(X_{old}, X_{new}) \\ \hline k(X_{new}, X_{old}) & k(X_{new}, X_{new}) + \sigma^2 I \end{array} =: \begin{array}{c|c} A & B \\ \hline C & D \end{array}.$$

If we have already conditioned on (X_{old}, X_{old}) , block A was already computed (see 1. above) and we only have to evaluate the remaining blocks B, C and D .

6. **predictive_mean**: The predictive mean also splits into two parts

$$m(X) + \mu(X) = \begin{pmatrix} m(X_{old}) + \mu(X_{old}) \\ m(X_{new}) + \mu(X_{new}) \end{pmatrix} =: \begin{pmatrix} a \\ b \end{pmatrix}.$$

Note that a was already computed, see 2. above.

7. **predictive_covariance_cho**: In order to compute a Cholesky factorization of $k(X, X) + \sigma^2 I$, we can use that we already have the Cholesky factor $A^{1/2}$ of A , see 3. above. It should be clear from theory exercise 1 (b) on the current exercise sheet how to construct the Cholesky factor for $k(X, X) + \sigma^2 I$.
8. **representer_weights**: We already computed the representer weights $A^{-1} \cdot u$, with $u = y_{old} - a$, see 4. above. Now, we have to solve a larger linear system

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix} \cdot \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} u \\ v \end{pmatrix}$$

with $v = y_{new} - b$ and the unknowns x and y . After computing v, y and x are given by

$$y = Q^{-1}(v - CA^{-1}u) \tag{1}$$

$$x = A^{-1}u - A^{-1}By, \tag{2}$$

where $Q = D - CA^{-1}B$ is the so-called **Schur-complement**. Note: All necessary quantities (like $Q^{1/2}$ and $A^{-1}u$) were already computed in previous steps.

Task: Implement an improved version of the condition method that constructs 5, 6, 7, and 8. Try to make this as efficient as possible by using the quantities 1, 2, 3, and 4. Make sure that the resulting posteriors are the same as the ones from approach 2 (e.g. by evaluating them on the same test data and comparing the Gaussians).

```
[ ]: from gaussians import Gaussian, ConditionalGaussianProcess

def condition_on_prior_improved(self, y_new, X_new, sigma):
    # Check data
    y_new = jnp.atleast_1d(y_new)
    X_new = jnp.atleast_2d(X_new)

    # Concatenate data
    X_both = jnp.concatenate([self.X, X_new], axis=0)
    y_both = jnp.concatenate([self.y, y_new], axis=0)

    # Concatenate Gaussians
    mu_new = jnp.zeros_like(y_new)
    mu_both = jnp.concatenate([self.epsilon.mu, mu_new], axis=0)

    Sigma_new = sigma * jnp.eye(len(y_new))
    Sigma_both = jax.scipy.linalg.block_diag(self.epsilon.Sigma, Sigma_new)

    epsilon_both = Gaussian(mu=mu_both, Sigma=Sigma_both)

    # Build conditional GP from prior (!)
    cond_gp = ConditionalGaussianProcess(self.prior, y=y_both, X=X_both,
                                          epsilon=epsilon_both)

    # =====
    # 5. predictive_covariance
    # =====
    # TODO
    A = self.predictive_covariance
    B = self.prior.k(self.X[:, None, :], X_new[None, :, :])
    D = self.prior.k(X_new[:, None, :], X_new[None, :, :]) + Sigma_new
    C = B.T

    def predictive_covariance_new(A, B, C, D):
        # Append A, B, C, D to a single square matrix of shape (n + m, n + m)
        # and return
        return jnp.block([[A, B], [C, D]])

    cond_gp.predictive_covariance = predictive_covariance_new(A, B, C, D)

    # =====
    # 6. predictive_mean
    # =====
    # TODO
    a = self.predictive_mean
```

```

b = self.prior.m(X_new) + mu_new
def predictive_mean_new(a, b):
    # Append a, b and return
    return jnp.concatenate([a, b], axis=0)

cond_gp.predictive_mean = predictive_mean_new(a, b)

# =====
# 7. predictive_covariance_cho
# =====
A_1_2, lower_A = self.predictive_covariance_cho
if not lower_A:
    # print('A is not lower triangular')
    A_1_2 = A_1_2.T
    lower_A = True
U_12 = jax.scipy.linalg.solve_triangular(A_1_2, B, lower=True)
L_21 = U_12.T
A_inverse_B = jax.scipy.linalg.cho_solve((A_1_2, lower_A), B)
Q = D - C @ A_inverse_B
Q_1_2, lower_Q = jax.scipy.linalg.cho_factor(Q, lower=True)
if not lower_Q:
    # print('Q is not lower triangular')
    Q_1_2 = Q_1_2.T
    lower_Q = True

def predictive_covariance_cho_alt(cond_gp):
    return jax.scipy.linalg.cho_factor(cond_gp.predictive_covariance)
def predictive_covariance_cho_new(A_1_2, U_12, L_21, Q_1_2):
    #U = jnp.block([[A_1_2.T, U_12], [jnp.zeros_like(L_21), Q_1_2.T]])
    L = jnp.block([[A_1_2, jnp.zeros_like(U_12)], [L_21, Q_1_2]])
    return (L, True)

#cond_gp.predictive_covariance_cho = predictive_covariance_cho_alt(cond_gp)
cond_gp.predictive_covariance_cho = predictive_covariance_cho_new(A_1_2, U_12, L_21, Q_1_2)

# =====
# 8. representer_weights
# =====
# TODO
u = self.y - a
v = y_new - b
repr_w_old = self.representer_weights
y_vec = jax.scipy.linalg.cho_solve((Q_1_2, lower_Q), (v-C @ repr_w_old))
x_vec = repr_w_old - A_inverse_B @ y_vec

```

```

def representer_weights_new(x_vec,y_vec):
    return jnp.concatenate([x_vec, y_vec], axis=0)

def representer_weights_alt():
    return jax.scipy.linalg.cho_solve(cond_gp.predictive_covariance_cho,
        ↪jnp.concatenate([u, v], axis=0))

#cond_gp.representer_weights = representer_weights_alt()
cond_gp.representer_weights = representer_weights_new(x_vec,y_vec)

return cond_gp

# Assign method to class
ConditionalGaussianProcess.condition = condition_on_prior_improved

# Evaluate approach 3
posteriors_A3 = get_posteriors()
eval_times_A3 = measure_eval_times()

# Make sure that the posteriors are identical to the ones from approach 2
# TODO
for posterior_A2, posterior_A3 in zip(posteriors_A2, posteriors_A3):
    same_mu = jnp.allclose(posterior_A2(X_test).mu, posterior_A3(X_test).mu)
    same_Sigma = jnp.allclose(posterior_A2(X_test).Sigma, posterior_A3(X_test).
        ↪Sigma)
    if not same_mu or not same_Sigma:
        print("The posteriors are not identical.")
        break
if same_mu and same_Sigma:
    print("The posteriors are identical.")

# Visualizations
plot_conditional_gps(posteriors_A3)
plot_eval_times([eval_times_A1, eval_times_A2, eval_times_A3])

```

The posteriors are identical.

