

Probabilistic Machine Learning

EBERHARD KARLS  
UNIVERSITÄT  
TÜBINGEN



**Exercise 12**

**Summer Term 2023**

Full Name	Matriculation No
Batuhan Özcömlekci	6300476
Aakarsh Nair	6546577

## Excercise-12

Batuhan, Oezcoemlekci (Matrikelnummer: 6300476)

Aakarsh, Nair (Matrikelnummer: 6546577)

16 June 2023

### Exercise

**Theory Question** Parameter inference with expectation maximization.

Consider a linear Gaussian model with the following structure:

- $x_0 \sim \mathcal{N}(m_0(\theta), P_0(\theta))$
- $x_k \sim \mathcal{N}(A(\theta)x_{k-1}, Q(\theta))$
- $y_k \sim \mathcal{N}(H(\theta)x_k, R(\theta))$

**To-do:** Estimate  $\theta$  for some data  $y_{1:N}$  an alternate to maximizing the likelihood (discussed in lecture) is the EM-Algorithm. It works as follows:

1. Start from an initial guess  $\theta^{(0)}$ .

2. For  $n = 0, 1, 2, \dots$  do:

(a) **E-Step:** Compute:

$$\mathcal{Q}(\theta, \theta^{(n)}) := \int p(x_{0:N} | y_{0:N}, \theta^{(n)}) \log p(x_{0:N}, y_{0:N} | \theta) dx_{0:N} \quad (1)$$

(b) **M-Step:** Compute:  $\theta^{(n+1)} := \arg \max_{\theta} \mathcal{Q}(\theta, \theta^{(n)})$

In a considered case of linear Gaussian state space model,  $\mathcal{Q}(\theta, \theta^{(n)})$  can be computed analytically. It is of the form:

$$\begin{aligned} \mathcal{Q}(\theta, \theta^{(n)}) = & -\frac{1}{2} \log |2\pi P_0(\theta)| - \frac{N}{2} \log |2\pi Q(\theta)| - \frac{N}{2} \log |2\pi R(\theta)| \\ & - \frac{1}{2} \text{tr} \left( P_0^{-1}(\theta) \left[ P_0^{(s)} + (m_0^{(s)} - m_0(\theta))(m_0^{(s)} - m_0(\theta))^T \right] \right) \\ & - \frac{1}{2} \text{tr} \left( Q^{-1}(\theta) \left[ \Sigma - CA(\theta)^T - A(\theta)C^T + A(\theta)\Phi A(\theta)^T \right] \right) \\ & - \frac{1}{2} \text{tr} \left( R^{-1}(\theta) \left[ D - BH(\theta)^T - H(\theta)B^T + H(\theta)\Sigma H(\theta)^T \right] \right) \end{aligned} \quad (2)$$

Where the following quantities are computed from the results of the Rauch-Tung-Striebel smoother run with parameter values  $\theta^{(n)}$ :

$$\Sigma = \frac{1}{N} \sum_{k=1}^N \left( P_k^{(s)} + m_k^s (m_k^s)^T \right)$$

$$\Phi = \frac{1}{N} \sum_{k=1}^N \left( P_{k-1}^{(s)} + m_{k-1}^s (m_{k-1}^s)^T \right)$$

$$B = \frac{1}{N} \sum_{k=1}^N \left( y_k (m_k^{(s)})^T \right)$$

$$C = \frac{1}{N} \sum_{k=1}^N \left( P_k^{(s)} G_{k-1}^T + m_k^{(s)} (m_{k-1}^s)^T \right)$$

$$D = \frac{1}{N} \sum_{k=1}^N \left( y_k y_k^T \right)$$

*Exercise:* Let  $\theta = (A, Q)$ , that is the model parameters are exactly the full transition model matrices. Derive closed form updates for both  $A$  and  $Q$  by computing the M-step.

(a) Compute  $A$ : We use the fact that trace is a linear operator.

$$\begin{aligned} \frac{\partial Q}{\partial A} &= 0 \\ \Rightarrow \frac{\partial}{\partial A} \left[ -\frac{N}{2} \text{tr} \left( Q^{-1} (\Sigma - CA^T - AC^T + A\Phi A^T) \right) \right] &= 0 \\ \Rightarrow \frac{N}{2} \frac{\partial}{\partial A} \text{tr} \left[ -Q^{-1} AC^T - Q^{-1} CA^T + Q^{-1} A\Phi A^T \right] &= 0 \\ \Rightarrow \frac{\partial}{\partial A} \text{tr} \left[ -Q^{-1} AC^T - Q^{-1} CA^T + Q^{-1} A\Phi A^T \right] &= 0 \\ \Rightarrow \frac{\partial}{\partial A} \text{tr} \left[ Q^{-1} A\Phi A^T \right] &= \frac{\partial}{\partial A} \text{tr} \left[ Q^{-1} AC^T + Q^{-1} CA^T \right] \end{aligned} \quad (3)$$

Now consider the LHS:

$$\frac{\partial}{\partial A} \text{tr} \left[ Q^{-1} A\Phi A^T \right] \quad (4)$$

This is known matrix differential form of the trace.

$$\frac{\partial}{\partial X_d} \text{tr} \left[ A_d X_d B_d X_d^T C_d \right] = A_d^T C_d^T X_d B_d^T + C_d A_d X_d B_d$$

Where  $A_d = Q^{-1}$ ,  $B_d = \Phi$ ,  $C_d = I$  and  $X_d = A$ .

Thus we have we use the symmetry of the postive semi-definite matrix inverse  $Q^{-1} = (Q^{-1})^T$  to simplify the LHS:

$$\begin{aligned} \frac{\partial}{\partial A} \text{tr} \left[ Q^{-1} A\Phi A^T \right] &= Q^{-T} I^T A\Phi^T + I Q^{-1} A\Phi \\ &= Q^{-T} A\Phi^T + Q^{-1} A\Phi \\ &= Q^{-T} A\Phi^T + Q^{-1} A\Phi \\ &= 2Q^{-1} A\Phi \end{aligned} \quad (5)$$

For the RHS we can we can simplify the first term using:

$$\frac{\partial}{\partial X_d} \text{tr} [A_d X_d B_d] = A_d^T B_d^T \quad (6)$$

With  $A_d = (Q^{-1})$  and  $X_d = A$  and  $B_d = C^T$  for the first term we get:

$$\frac{\partial}{\partial A} \text{tr} [Q^{-1} A C^T] = Q^{-T} C = Q^{-1} C \quad (7)$$

And for the second term we can use the following matrix identity we get:

$$\frac{\partial}{\partial X_d} \text{tr} [A_d X_d^T] = A_d \quad (8)$$

With  $A_d = (Q^{-1})C$  and  $X_d = A$  the second term simplifies to:

$$\frac{\partial}{\partial A} \text{tr} [Q^{-1} C A^T] = Q^{-1} C \quad (9)$$

Thus taken together , using the linearity of the trace operator we can separate the terms we get:

$$\begin{aligned} \frac{\partial}{\partial A} \text{tr} [Q^{-1} A C^T + Q^{-1} C A^T] &= Q^{-1} C + Q^{-1} C \\ &= 2Q^{-1} C \end{aligned} \quad (10)$$

Thus LHS and RHS together becomes:

$$\begin{aligned} 2Q^{-1} A \Phi &= 2Q^{-1} C \\ Q^{-1} A \Phi &= Q^{-1} C \\ A \Phi &= C \\ A &= C \Phi^{-1} \end{aligned} \quad (11)$$

Thus we have  $A^* = C \Phi^{-1}$ .

(b) Compute  $Q$ :

We again compute the critical point by computing the matrix derivative with respect to  $Q$  and setting it to zero. Canceling the the constant from the linear operators for trace and derivative.

Let

$$Z = [\Sigma - C A^T - A C^T + A \Phi A^T]$$

$$\begin{aligned}
\frac{\partial}{\partial Q} \left[ -\frac{N}{2} \log |2\pi Q| - \frac{N}{2} \text{tr} (Q^{-1}Z) \right] &= 0 \\
\frac{\partial}{\partial Q} \left[ -\log |2\pi Q| - \frac{\partial}{\partial Q} \text{tr} (Q^{-1}Z) \right] &= 0 \\
\frac{\partial}{\partial Q} [-\log |2\pi Q|] - \left[ \frac{\partial}{\partial Q} \text{tr} (Q^{-1}Z) \right] &= 0 \quad (12) \\
\frac{\partial}{\partial Q} [-\log |Q|] - \left[ \frac{\partial}{\partial Q} \text{tr} (Q^{-1}Z) \right] &= 0 \\
- \left[ \frac{\partial}{\partial Q} \text{tr} (Q^{-1}Z) \right] &= \frac{\partial}{\partial Q} [\log |Q|]
\end{aligned}$$

Consider the LHS, we use the following matrix differential form of the trace.

$$\frac{\partial}{\partial X_d} \text{tr}(A_d X_d^{-1} B_d) = -X_d^{-T} A_d^T B_d^T X_d^{-T} \quad (13)$$

With  $A_d = I$ ,  $B_d = Z$  and  $X_d = Q$ , and the fact that  $Q^{-T} = Q^{-1}$  being inverse of PSD matrix we get:

$$\frac{\partial}{\partial Q} (Q^{-1}Z) = -Q^{-T} Z^T Q^{-T} = -Q^{-1} Z^T Q^{-1} \quad (14)$$

We also note that  $Z = Z^T$  as we have that  $\Sigma$  and  $\Phi$  are symmetric matrices.

$$\begin{aligned}
Z^T &= (\Sigma - CA^T - AC^T + A\Phi A^T)^T = (\Sigma^T - AC^T - CA^T + A\Phi^T A^T) \\
&= (\Sigma - CA^T - AC^T + A\Phi A^T) = Z
\end{aligned} \quad (15)$$

Thus we get the LHS to be :

$$\frac{\partial}{\partial X_d} \text{tr}(Q^{-1}Z) = -Q^{-1}ZQ^{-1} \quad (16)$$

For the RHS we use the matrix identity:

$$\frac{\partial}{\partial X_d} \log |X_d| = X_d^{-T} \quad (17)$$

With  $X_d = Q$  we get the RHS to be:

$$\frac{\partial}{\partial Q} \log |Q| = Q^{-T} = Q^{-1} \quad (18)$$

Taken together we get

$$\begin{aligned}
 -Q^{-1}ZQ^{-1} &= -Q^{-1} \\
 \implies Q^{-1}ZQ^{-1} &= Q^{-1} \quad (19) \\
 \implies ZQ^{-1} &= I
 \end{aligned}$$

Thus from the uniqueness of the inverse we get:

$$Q^* = Z = (\Sigma - CA^T - AC^T + A\Phi A^T) \quad (20)$$

### *References*

# Ex12\_code

July 24, 2023

## 1 Probabilistic Machine Learning

University of Tübingen, Summer Term 2023 © 2023 P. Hennig

### 1.1 Exercise Sheet No. 12

---

Submission by:

- Batuhan, Oezcoemlekci, Matrikelnummer: 6300476
- Aakarsh, Nair, Matrikelnummer: 6546577

```
[ ]: import numpy as np

from matplotlib import pyplot as plt
from numpy.typing import ArrayLike

import scipy.io as sio
import scipy.linalg as sla
import scipy.special as ssp

from tueplots import bundles
from tueplots.constants.color import rgb

# plt.rcParams.update(bundles.beamer_moml())
plt.rcParams.update({"figure.dpi": 200})
```

In this exercise, you will implement a Rauch-Tung-Striebel (RTS) smoother. We will work on the same data as before so you will recognize parts of the notebook from the previous exercise.

As before, only change code in cells where we explicitly ask you to.

## 2 I. The Data

```
[ ]: DIM = 7
NUM_DERIV = 2
STATE_DIM = DIM * (NUM_DERIV + 1)
```

```
[ ]: proj_position = np.eye(STATE_DIM)[:DIM, :]
proj_velocity = np.eye(STATE_DIM)[DIM:2*DIM, :]
proj_acceleration = np.eye(STATE_DIM)[2*DIM:, :]
```

```
[ ]: def plot_data(axes, Y):
    assert len(axes) == 3
    N, d = Y.shape
    num_joints = d // 3
    xs = np.arange(N)
    positions = Y @ proj_position.T
    velocities = Y @ proj_velocity.T
    accelerations = Y @ proj_acceleration.T
    for i in range(num_joints):
        axes[0].scatter(xs, positions[:, i], marker="x", s=5, label="joint {}".
        ↪format(i), color="C{}".format(i))
        axes[1].scatter(xs, velocities[:, i], marker="x", s=5, label="joint {}".
        ↪format(i), color="C{}".format(i))
        axes[2].scatter(xs, accelerations[:, i], marker="x", s=5, label="joint_
        ↪{}".format(i), color="C{}".format(i))
    # plt.legend()
    return axes
```

```
[ ]: def plot_estimate(axes, kf_means, kf_covs, fctr=1.97):
    assert len(axes) == 3
    N, d = kf_means.shape
    num_joints = d // 3
    xs = np.arange(N)
    m_positions = kf_means @ proj_position.T
    m_velocities = kf_means @ proj_velocity.T
    m_accelerations = kf_means @ proj_acceleration.T

    kf_stds = np.array([fctr * np.sqrt(np.diag(C)) for C in kf_covs])
    s_positions = kf_stds @ proj_position.T
    s_velocities = kf_stds @ proj_velocity.T
    s_accelerations = kf_stds @ proj_acceleration.T

    for i in range(num_joints):
        axes[0].plot(xs, m_positions[:, i], color="C{}".format(i))
        axes[0].fill_between(xs, m_positions[:, i] - s_positions[:, i],
        ↪m_positions[:, i] + s_positions[:, i], color="C{}".format(i), alpha=0.4)
        axes[1].plot(xs, m_velocities[:, i], color="C{}".format(i))
        axes[1].fill_between(xs, m_velocities[:, i] - s_velocities[:, i],
        ↪m_velocities[:, i] + s_velocities[:, i], color="C{}".format(i), alpha=0.4)
        axes[2].plot(xs, m_accelerations[:, i], color="C{}".format(i))
        axes[2].fill_between(xs, m_accelerations[:, i] - s_accelerations[:, i],
        ↪m_accelerations[:, i] + s_accelerations[:, i], color="C{}".format(i),
        ↪alpha=0.4)
```



```
# plt.legend()
return axs
```

```
[ ]: data = sio.loadmat('sarcos_inv.mat')['sarcos_inv'][:, :-7]
Y = data
```

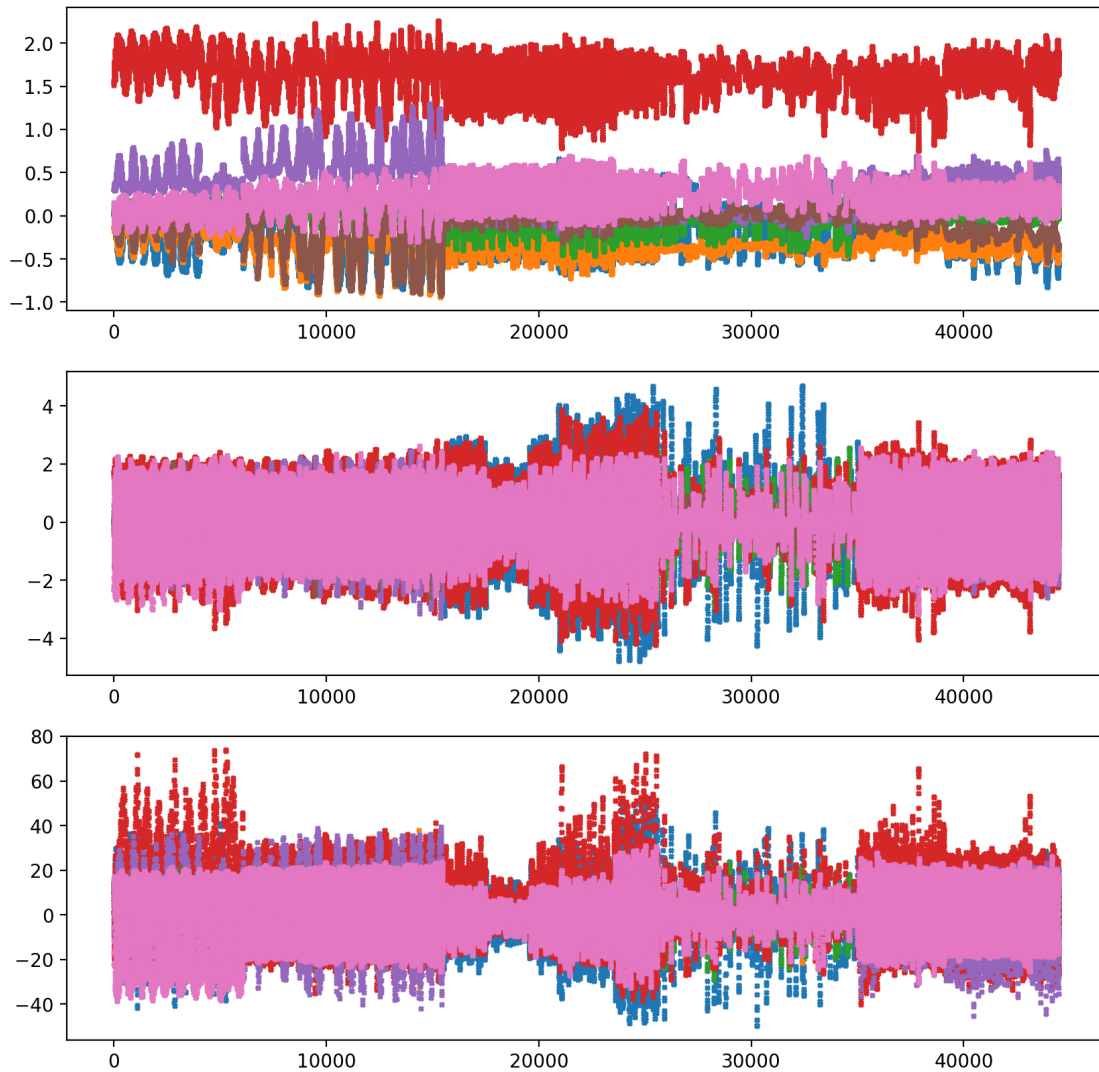
```
[ ]: data.shape
```

```
[ ]: (44484, 21)
```

**2.1** Now, let's have a look at the entire time series.

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
plot_data(axs, data)
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```




Pretty chaotic, huh... Well, it's over 44 thousand data points and our screen is only so wide... That's why we are going to look at a smaller, zoomed-in window from now on.

```
[ ]: time_window_for_viz = slice(11000, 11050)
```

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
      plot_data(axs, data[time_window_for_viz, :])
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```



### 3 II. The Model

Ok, now that we have a feel for the data and what it looks like, we are going to set up a model. What kind of model this is, we are going to be secretive about for now. Perhaps, you will learn about it in one of the following lectures? (Perhaps not, let's see).

**The following two cells ...** ... are mysterious functions that create us two matrices  $A$  and  $Q$ , the transition matrix and process-noise covariance matrix of our linear, time-invariant Gaussian transition density.

**You do not have to understand what these two functions do, just take them for granted!** (I wouldn't try, anyway...)

**The dynamics model (prior)**

```
[ ]: def mysterious_operation(F, L, dt):
    dim = F.shape[0]

    if L.ndim == 1:
        L = L.reshape((-1, 1))

    Phi = np.block(
        [
            [F, L @ L.T],
            [np.zeros(F.shape), -F.T],
        ]
    )
    M = sla.expm(Phi * dt)

    Ah = M[:dim, :dim]
    Qh = M[:dim, dim:] @ Ah.T

    return Ah, Qh
```

```
[ ]: def create_mysterious_ssm(d, q, ell, dt, dff=1.0):
    F = np.diag(np.ones(q), 1)
    nu = q + 0.5
    D, lam = q + 1, np.sqrt(2 * nu) / ell
    F[-1, :] = np.array(
        [-ssp.binom(D, i) * lam ** (D - i) for i in range(D)]
    )

    L = np.eye(q+1)[-1, :] * dff
    A_1d, Q_1d = mysterious_operation(F, L, dt)
    A = np.kron(A_1d, np.eye(d))
    Q = np.kron(Q_1d, np.eye(d))
    return A, Q
```

```
[ ]: dt = 1.0
```

```
[ ]: A, Q = create_mysterious_ssm(DIM, NUM_DERIV, 10.0, dt, 50.0)
```

### The measurement model (likelihood)

```
[ ]: H = np.eye(STATE_DIM) # We measure the entire state.
R = np.kron(np.diag(np.array([0.01, 0.01, 1.0])), np.eye(DIM)) # a (not-quite-
↪isotropic) sensor noise.
```

Finally, we set the initial moments to the first data point with some uncertainty.

```
[ ]: m0 = data[0, :]
P0 = np.kron(np.diag(np.array([0.01, 0.01, 1.0])), np.eye(DIM))
```

## 4 III. Inference

### 4.1 Step 1: Filtering

```
[ ]: def symmetrize(A):  
    return 0.5 * (A + A.T) + (1e-8 * np.eye(A.shape[0]))
```

```
[ ]: def kf_predict(m_filt, P_filt, A, Q):  
    m_pred = A @ m_filt  
    P_pred = A @ P_filt @ A.T + Q  
    return m_pred, symmetrize(P_pred)
```

```
[ ]: def kf_update(m_pred, P_pred, H, R, y):  
    predicted_measurement = H @ m_pred  
    innovation = (y - predicted_measurement)  
    innovation_gramian = H @ P_pred @ H.T + R  
  
    S_chol_fact = sla.cho_factor(symmetrize(innovation_gramian))  
    cross_covariance = P_pred @ H.T  
    kalman_gain = sla.cho_solve(S_chol_fact, cross_covariance.T).T  
  
    mean_increment = kalman_gain @ innovation  
    covariance_decrement = kalman_gain @ innovation_gramian @ kalman_gain.T  
  
    m_filt = m_pred + mean_increment  
    P_filt = P_pred - covariance_decrement  
    return m_filt, symmetrize(P_filt)
```


```
[ ]: def filter_kalman(m0, P0, A, Q, H, R, Y):  
    d, D = H.shape  
    N = Y.shape[0]  
    result_mean = [m0.copy()]  
    result_cov = [P0.copy()]  
    m = m0.copy()  
    P = P0.copy()  
    for n in range(1, N):  
        m, P = kf_predict(m, P, A, Q)  
        m, P = kf_update(m, P, H, R, Y[n, :])  
        result_mean.append(m.copy())  
        result_cov.append(P.copy())  
    return np.array(result_mean), np.array(result_cov)
```

```
[ ]: %%time  
kf_means, kf_covs = filter_kalman(m0, P0, A, Q, H, R, Y)
```

CPU times: user 5.29 s, sys: 72 ms, total: 5.36 s  
Wall time: 5.36 s

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
plot_estimate(axs, kf_means[time_window_for_viz, :],
             kf_covs[time_window_for_viz, :, :])
plot_data(axs, data[time_window_for_viz, :])
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```



## 4.2 Step 2: Smoothing

### 4.3 Task 1:

Implement a Rauch-Tung-Striebel smoother. ### (a) Fill the function body of the function `smoother_step` below. It takes as arguments - `filt_m`, `filt_P`: the filtering moments at time step  $k-1$  - `pred_m`, `pred_P`: the predicted moments at time step  $k$  - `smooth_m`, `smooth_P`: the smoothed

moments at time step  $k$  - A, Q: the parameters of the transition density

The function must return a tuple containing 1. xi the smoothing mean at time step  $k-1$  2. Lambda the smoothing covariance at time step  $k-1$  3. G the smoothing gain used to compute the above

```
[ ]: def smoother_step(filt_m, filt_P, pred_m, pred_P, smooth_m, smooth_P, A, Q):
    pred_P_cho_factor, _ = sla.cho_factor(pred_P, lower=True)
    pred_P_inverse = sla.cho_solve((pred_P_cho_factor, True), np.eye(pred_P.
    ↪shape[0]))
    #pred_P_inverse = sla.inv(pred_P_inverse)
    G = filt_P @ A.T @ pred_P_inverse
    xi = filt_m + G @ (smooth_m - pred_m)
    Lambda = filt_P + G @ (smooth_P - pred_P) @ G.T
    return xi, Lambda, G
```

#### 4.4 (b)

Fill the function body of the function `rts_smooth` below. It takes as arguments - `filter_means` - `filter_covs` - A, Q: the parameters of the transition density

##### 4.4.1 IMPORTANT

This function must return a tuple of three arrays: - an  $N \times D$ -array containing the smoother means - an  $N \times D \times D$ -array containing the smoother covariances - AND an  $N \times D \times D$ -array containing the smoother gains  $G_k$  from every step. We will need the last one for later.

```
[ ]: def rts_smooth(filter_means, filter_covs, A, Q):
    smoother_means = np.zeros_like(filter_means)
    smoother_covs = np.zeros_like(filter_covs)
    smoother_gains = np.zeros_like(filter_covs)

    for i in range(filter_means.shape[0], 0, -1):
        if i == filter_means.shape[0]:
            smoother_means[i-1, :] = filter_means[i-1, :]
            smoother_covs[i-1, :, :] = filter_covs[i-1, :, :]
        else:
            smoother_means[i-1, :], smoother_covs[i-1, :, :], ↪
    ↪smoother_gains[i-1, :, :] = smoother_step(
                filter_means[i-1, :], filter_covs[i-1, :, :],
                filter_means[i, :], filter_covs[i, :, :],
                smoother_means[i, :], smoother_covs[i, :, :],
                A, Q
            )

    return smoother_means, smoother_covs, smoother_gains
```

## 5 Now we test your implementation.

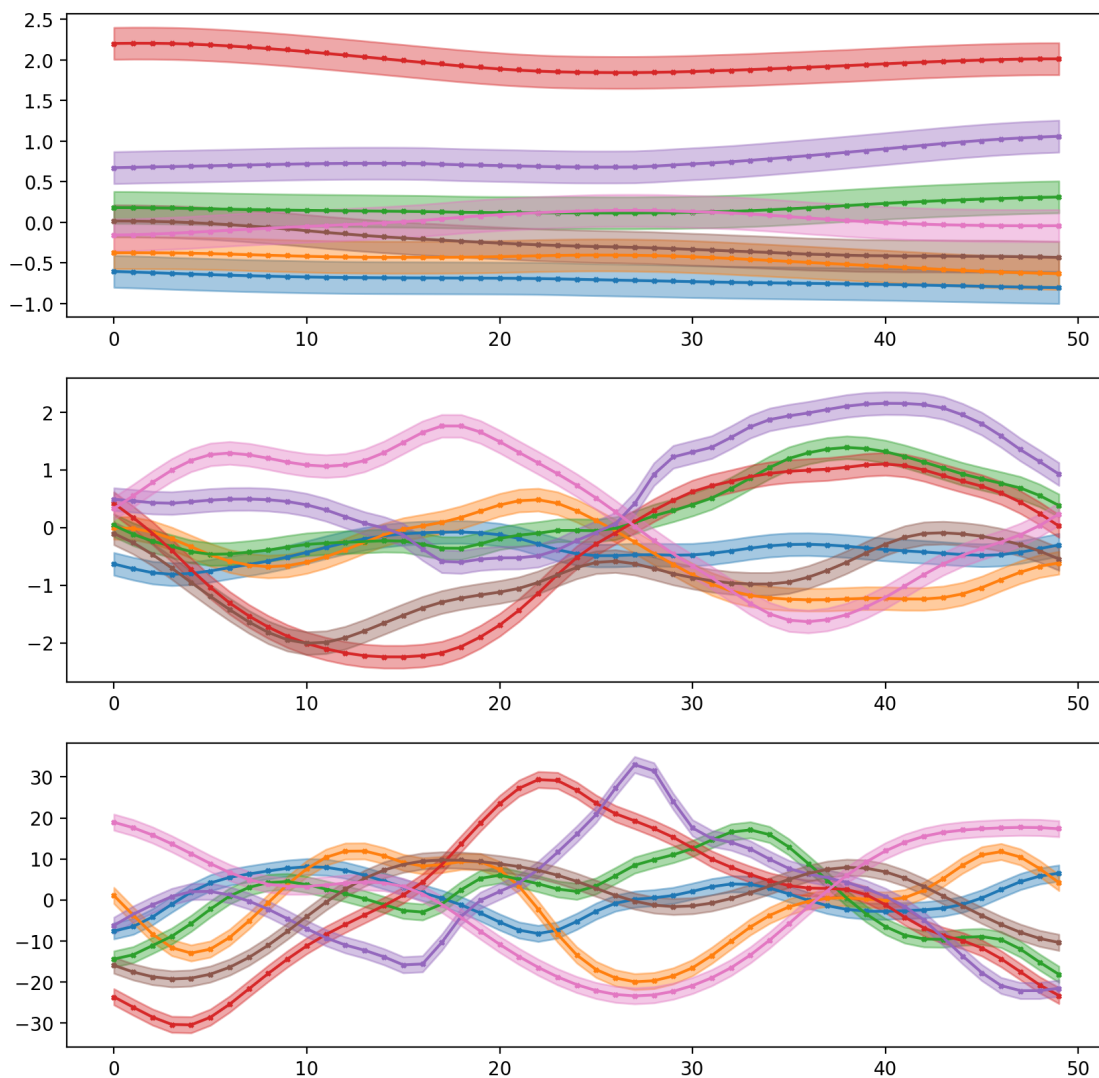
DO NOT CHANGE ANYTHING IN THESE CELLS

```
[ ]: %%time
rts_means, rts_covs, _ = rts_smooth(kf_means, kf_covs, A, Q)
```

CPU times: user 2.76 s, sys: 60.2 ms, total: 2.83 s  
Wall time: 2.84 s

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
plot_estimate(axs, rts_means[time_window_for_viz, :],
             ↪rts_covs[time_window_for_viz, :, :])
plot_data(axs, data[time_window_for_viz, :])
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```





---


## 5.1 Task 2

All right, now - say - we do not have a good model for the data we see. Take for example this transition model here, which is really just random Gaussian white noise in every step.

```
[ ]: A_init = np.zeros((STATE_DIM, STATE_DIM))
     Q_init = 0.01*np.diag(np.random.rand(STATE_DIM))
```

```
[ ]: fig, axs = plt.subplots(1, 2, figsize=(10, 10))
     axs[0].imshow(A_init)
     axs[1].imshow(Q_init)
```


```
[ ]: <matplotlib.image.AxesImage at 0x7f0b86305490>
```



```
[ ]: init_kf_means, init_kf_covs = filter_kalman(m0, P0, A_init, Q_init, H, R, Y)
```

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
     plot_estimate(axs, init_kf_means[time_window_for_viz, :],
     ↪init_kf_covs[time_window_for_viz, :, :])
     plot_data(axs, data[time_window_for_viz, :])
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```



Surprise, the fit is not so good. Luckily, you know a tool that might help.

## 6 The EM-Algorithm for linear Gaussian state-space models

### 6.0.1 (a) Implement the E-step of the EM algorithm

as given in the theory exercise. The function takes -  $m_0$ ,  $P_0$ : initial moments -  $A$ ,  $Q$ : a transition model -  $H$ ,  $R$ : a measurement model -  $Y$

The E-step computes a bunch of matrices

- $\Sigma$
- $\Phi$
- $B$
- $C$

- D

from the RTS smoother estimate given the current SSM and returns them.

```
[ ]: def E_step(m0, P0, A, Q, H, R, Y):
    N = Y.shape[0]
    _kf_means, _kf_covs = filter_kalman(m0, P0, A, Q, H, R, Y)

    for i in range(_kf_covs.shape[0] - 1, 0, -1):
        assert np.all(np.linalg.eigvals(_kf_covs[i, :, :]) > 0), f"matrix is not psd {np.linalg.eigvals(_kf_covs[i, :, :])}"
        assert np.allclose(_kf_covs[i, :, :], _kf_covs[i, :, :].T), "covariance is not symmetric"

    rts_means, rts_covs, rts_gains = rts_smooth(_kf_means, _kf_covs, A, Q)

    Sigma = (1.0 / N) * np.sum([rts_covs[k, :, :] + np.outer(rts_means[k, :] ,
    ↪ rts_means[k, :].T) for k in range(N)], axis=0)
    Phi = (1.0 / N) * np.sum([rts_covs[k-1, :, :] + np.outer(rts_means[k-1, :]
    ↪ , rts_means[k-1, :].T) for k in range(1, N)], axis=0)
    B = (1.0 / N) * np.sum([np.outer(Y[k, :] , rts_means[k, :].T) for k in
    ↪ range(N)], axis=0)
    # C - uses previous gains unlike term in assignment.
    C = (1.0 / N) * np.sum([ rts_covs[k, :, :] @ rts_gains[k-1, :, :].T + np.
    ↪ outer(rts_means[k, :] , rts_means[k-1, :].T) for k in range(N)], axis=0)
    D = (1.0 / N) * np.sum([np.outer(Y[k, :] , Y[k, :]) for k in range(N)],
    ↪ axis=0)

    assert Sigma.shape == (STATE_DIM, STATE_DIM)
    assert Phi.shape == (STATE_DIM, STATE_DIM)
    assert B.shape == (STATE_DIM, STATE_DIM)
    assert C.shape == (STATE_DIM, STATE_DIM)
    assert D.shape == (STATE_DIM, STATE_DIM)

    return symmetrize(Sigma), symmetrize(Phi), B, C, D
```

## 6.0.2 (b) Compute the M-Step for the transition matrix A

based on the results of the E-Step

```
[ ]: def M_step_A(Sigma, Phi, B, C, D):
    # Your code goes here.
    Phi_cho_factor, _ = sla.cho_factor(symmetrize(Phi), lower=True)
    Phi_inv = sla.cho_solve((Phi_cho_factor, True), np.eye(Phi.shape[0]))
```

```

    assert np.allclose(symmetrize(Phi) @ Phi_inv, np.eye(Phi.shape[0]),
        rtol=1e-4, atol=1e-6), "inverse is wrong, {}".format(np.linalg.norm(Phi @
        Phi_inv - np.eye(Phi.shape[0])))
    return C @ Phi_inv

```

### 6.0.3 (c) Compute the M-Step for the transition noise covariance $Q$

based on the results of the E-Step

```

[ ]: def M_step_Q(Sigma, Phi, B, C, D, A):
    # Your code goes here.
    return symmetrize(Sigma - C @ A.T - A @ C.T + A @ Phi @ A.T)

```

## 7 From here on, DON'T CHANGE ANYTHING.

### 7.0.1 This might take a while

```

[ ]: def EM_AQ(m0, P0, A_init, Q_init, H, R, Y, n_iter):
    A_star = A_init.copy()
    Q_star = Q_init.copy()
    for i in range(n_iter):
        print("EM-Step {}".format(i+1))
        Sigma, Phi, B, C, D = E_step(m0, P0, A, Q_star, H, R, Y)
        A_star = M_step_A(Sigma, Phi, B, C, D)
        Q_star = M_step_Q(Sigma, Phi, B, C, D, A)
    return A_star, Q_star

```

```

[ ]: A_star, Q_star = EM_AQ(m0, P0, A_init, Q_init, H, R, Y, 100)

```

```

EM-Step 1
EM-Step 2
EM-Step 3
EM-Step 4
EM-Step 5
EM-Step 6
EM-Step 7
EM-Step 8
EM-Step 9
EM-Step 10
EM-Step 11
EM-Step 12
EM-Step 13
EM-Step 14
EM-Step 15
EM-Step 16
EM-Step 17


```

EM-Step 18  
EM-Step 19  
EM-Step 20  
EM-Step 21  
EM-Step 22  
EM-Step 23  
EM-Step 24  
EM-Step 25  
EM-Step 26  
EM-Step 27  
EM-Step 28  
EM-Step 29  
EM-Step 30  
EM-Step 31  
EM-Step 32  
EM-Step 33  
EM-Step 34  
EM-Step 35  
EM-Step 36  
EM-Step 37  
EM-Step 38  
EM-Step 39  
EM-Step 40  
EM-Step 41  
EM-Step 42  
EM-Step 43  
EM-Step 44  
EM-Step 45  
EM-Step 46  
EM-Step 47  
EM-Step 48  
EM-Step 49  
EM-Step 50  
EM-Step 51  
EM-Step 52  
EM-Step 53  
EM-Step 54  
EM-Step 55  
EM-Step 56  
EM-Step 57  
EM-Step 58  
EM-Step 59  
EM-Step 60  
EM-Step 61  
EM-Step 62  
EM-Step 63  
EM-Step 64  
EM-Step 65

EM-Step 66  
EM-Step 67  
EM-Step 68  
EM-Step 69  
EM-Step 70  
EM-Step 71  
EM-Step 72  
EM-Step 73  
EM-Step 74  
EM-Step 75  
EM-Step 76  
EM-Step 77  
EM-Step 78  
EM-Step 79  
EM-Step 80  
EM-Step 81  
EM-Step 82  
EM-Step 83  
EM-Step 84  
EM-Step 85  
EM-Step 86  
EM-Step 87  
EM-Step 88  
EM-Step 89  
EM-Step 90  
EM-Step 91  
EM-Step 92  
EM-Step 93  
EM-Step 94  
EM-Step 95  
EM-Step 96  
EM-Step 97  
EM-Step 98  
EM-Step 99  
EM-Step 100

```
[ ]: fig, axs = plt.subplots(1, 2, figsize=(10, 10))  
    axs[0].imshow(A_star)  
    axs[1].imshow(Q_star)
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f0b858dd490>
```




```
[ ]: %%time
kf_means, kf_covs = filter_kalman(m0, P0, A_star, Q_star, H, R, Y)
```

CPU times: user 4.71 s, sys: 60 ms, total: 4.77 s  
Wall time: 4.77 s

```
[ ]: fig, axs = plt.subplots(3,1, figsize=(10, 10))
plot_estimate(axs, kf_means[time_window_for_viz, :],
↳kf_covs[time_window_for_viz, :, :])
plot_data(axs, data[time_window_for_viz, :])
```

```
[ ]: array([<Axes: >, <Axes: >, <Axes: >], dtype=object)
```



## 7.1 If everything went as planned, ...

... you should see a good model with a fine fit now.

### 7.1.1 How to submit your work:

Export your answer into a pdf (for example using jupyter's **Save** and **Export Notebook** as feature in the **File** menu). Make sure to include all outputs, in particular plots. Also include your answer to the theory question, either by adding it as LaTeX code directly in the notebook, or by adding it as an extra page (e.g. a scan) to the pdf. Submit the exercise on Ilias, in the associated folder. **Do not forget to add your name(s) and matrikel number(s) above!**